GPT-OSS-20B — actual evidence-coverage distributions at completion
Gray (up) = Clean; color (down, mirrored) = the labeled condition — same paired examples, no derived difference

Epistemic Confirmation
mean cov.: Clean=82.6%, cond.=83.1% (n=1193)

Search Match
mean cov.: Clean=82.6%, cond.=82.2% (n=1169)

Source Coverage
mean cov.: Clean=82.3%, cond.=82.0% (n=1189)

Epistemic Confirmation + Source Coverage
mean cov.: Clean=82.6%, cond.=83.0% (n=1190)

Search Match + Source Coverage
mean cov.: Clean=82.3%, cond.=81.8% (n=1172)

Fixed Sufficiency
mean cov.: Clean=82.6%, cond.=81.1% (n=1199)

Fixed Sufficiency Saturation
mean cov.: Clean=82.4%, cond.=81.4% (n=1200)

Oracle Complete--Neutral
mean cov.: Clean=82.8%, cond.=100.0% (n=1204)

Oracle Complete--Saturation
mean cov.: Clean=82.6%, cond.=100.0% (n=1194)

% of examples
(up = Clean, down = condition)

50 25 0 25 50

0 25 33.3 50 66.7 75 100

Evidence coverage at completion (%) — actual value, not a delta